

Lab 2 Result Report:

Table 1: Find Max Number of Connections:

Concurrency Level	Number of Requests	Requests Handled
100	100	100
200	200	200
300	300	270
270	270	259
250	250	248
240	240	234
210	210	204

From the table above, we can conclude that the **max number** of requests that can be handled before any connection drops is : **200**.

Terminal Result:

```
ubuntu@ip-172-31-42-27:~$ ab -n 200 -c 200
http://ec2-34-235-2-173.compute-1.amazonaws.com/?keywords=helloworld+foo+bar
This is ApacheBench, Version 2.3 <$Revision: 1528965 $>
Copyright 1996 Adam Twiss, Zeus Technology Ltd, http://www.zeustech.net/
Licensed to The Apache Software Foundation, http://www.apache.org/
```

Benchmarking ec2-34-235-2-173.compute-1.amazonaws.com (be patient)

Completed 100 requests

apr_socket_recv: Connection reset by peer (104)

Total of 116 requests completed

```
ubuntu@ip-172-31-42-27:~$ ab -n 200 -c 200
```

```
http://ec2-34-235-2-173.compute-1.amazonaws.com/?keywords=helloworld+foo+bar
```

This is ApacheBench, Version 2.3 <\$Revision: 1528965 \$>

Copyright 1996 Adam Twiss, Zeus Technology Ltd, http://www.zeustech.net/

Licensed to The Apache Software Foundation, http://www.apache.org/

Benchmarking ec2-34-235-2-173.compute-1.amazonaws.com (be patient)

Completed 100 requests

Completed 200 requests

Finished 200 requests

Server Software: WSGIServer/0.1
Server Hostname: ec2-34-235-2-173.compute-1.amazonaws.com
Server Port: 80

Document Path: /?keywords=helloworld+foo+bar
Document Length: 1919 bytes

Concurrency Level: 200
Time taken for tests: 0.856 seconds
Complete requests: 200
Failed requests: 0
Total transferred: 437200 bytes
HTML transferred: 383800 bytes
Requests per second: 233.52 [#/sec] (mean)
Time per request: 856.449 [ms] (mean)
Time per request: 4.282 [ms] (mean, across all concurrent requests)
Transfer rate: 498.52 [Kbytes/sec] received

Connection Times (ms)

	min	mean[+/-sd]	median	max
Connect:	0	1 1.8	0	8
Processing:	5	47 127.5	27	848
Waiting:	4	46 127.5	26	848
Total:	13	48 128.7	27	856

Percentage of the requests served within a certain time (ms)

50%	27
66%	28
75%	28
80%	29
90%	30
95%	30
98%	845
99%	856
100%	856 (longest request)

From the table and benchmark output above, we get the following conclusions: (see highlighted content in red)

- Max number of requests that can be handled: 200
- Max Number of Requests Per Second (RPS): 233.52 [#/sec] (mean)
- Average Response Time: 27ms
99 Percentile Response Time: 856ms